GRANTBRIDGE · GRANT STRATEGY FOR SMALL NONPROFITS

Less time chasing funding. More time changing lives.

Most grant tools help you apply. GrantBridge helps a small nonprofit decide **what is worth applying for**, and **whether it can afford to win**.



Photo: Sweet Life on Unsplash (Unsplash License). The people shown are not affiliated with GrantBridge.

THE PROBLEM

Small teams lose the most before they write a word

92%

of U.S. nonprofits spend less than $1 million a year.

45%

of nonprofits with government grants or contracts reported problems with late payments, in a national survey (2012 data)

- Applications compete with programs for the same few staff hours.
- The costly mistakes come *before* writing: an ineligible grant, colliding deadlines, mission drift.
- And winning can cost cash first: a required match, and payment only after you spend.
- Writing assistants start at the draft. GrantBridge starts at the decision.

Sources: National Council of Nonprofits, *Nonprofit Impact Matters* (2019 report); Urban Institute, *Nonprofit-Government Contracts and Grants: Findings from the 2013 National Survey*, p. 36. Both accessed 2026-10-21.

WHAT IS NEW

Everything before the writing, and everything after it

Discovery tools end where the hard part begins; writing assistants start after the expensive decision is already made. GrantBridge is the five steps in between — and code, not the model, makes every call that has to be right.

1 · Should we pursue it?

Eligibility, fit, readiness, staff hours and budget, each from a rule printed on the page. The AI adds evidence; it cannot override a check.

No fortnight spent on a grant you were never eligible for.

2 · Can we afford to win it? NEW

Quotes the match, the reimbursement delay and the overhead cap, then simulates the cash winning would tie up.

Catches the grant that fits but drains the bank account.

3 · Which ones, together?

Grants scheduled against the team's real weekly hours; mission drift flagged in the funder's own words.

A plan the team can actually staff.

4 · The funder's own outline NEW

The solicitation read into the sections it prescribes, what each must contain, the forms and the rubric.

The hour of copying that starts every application.

5 · Did we actually answer it? NEW

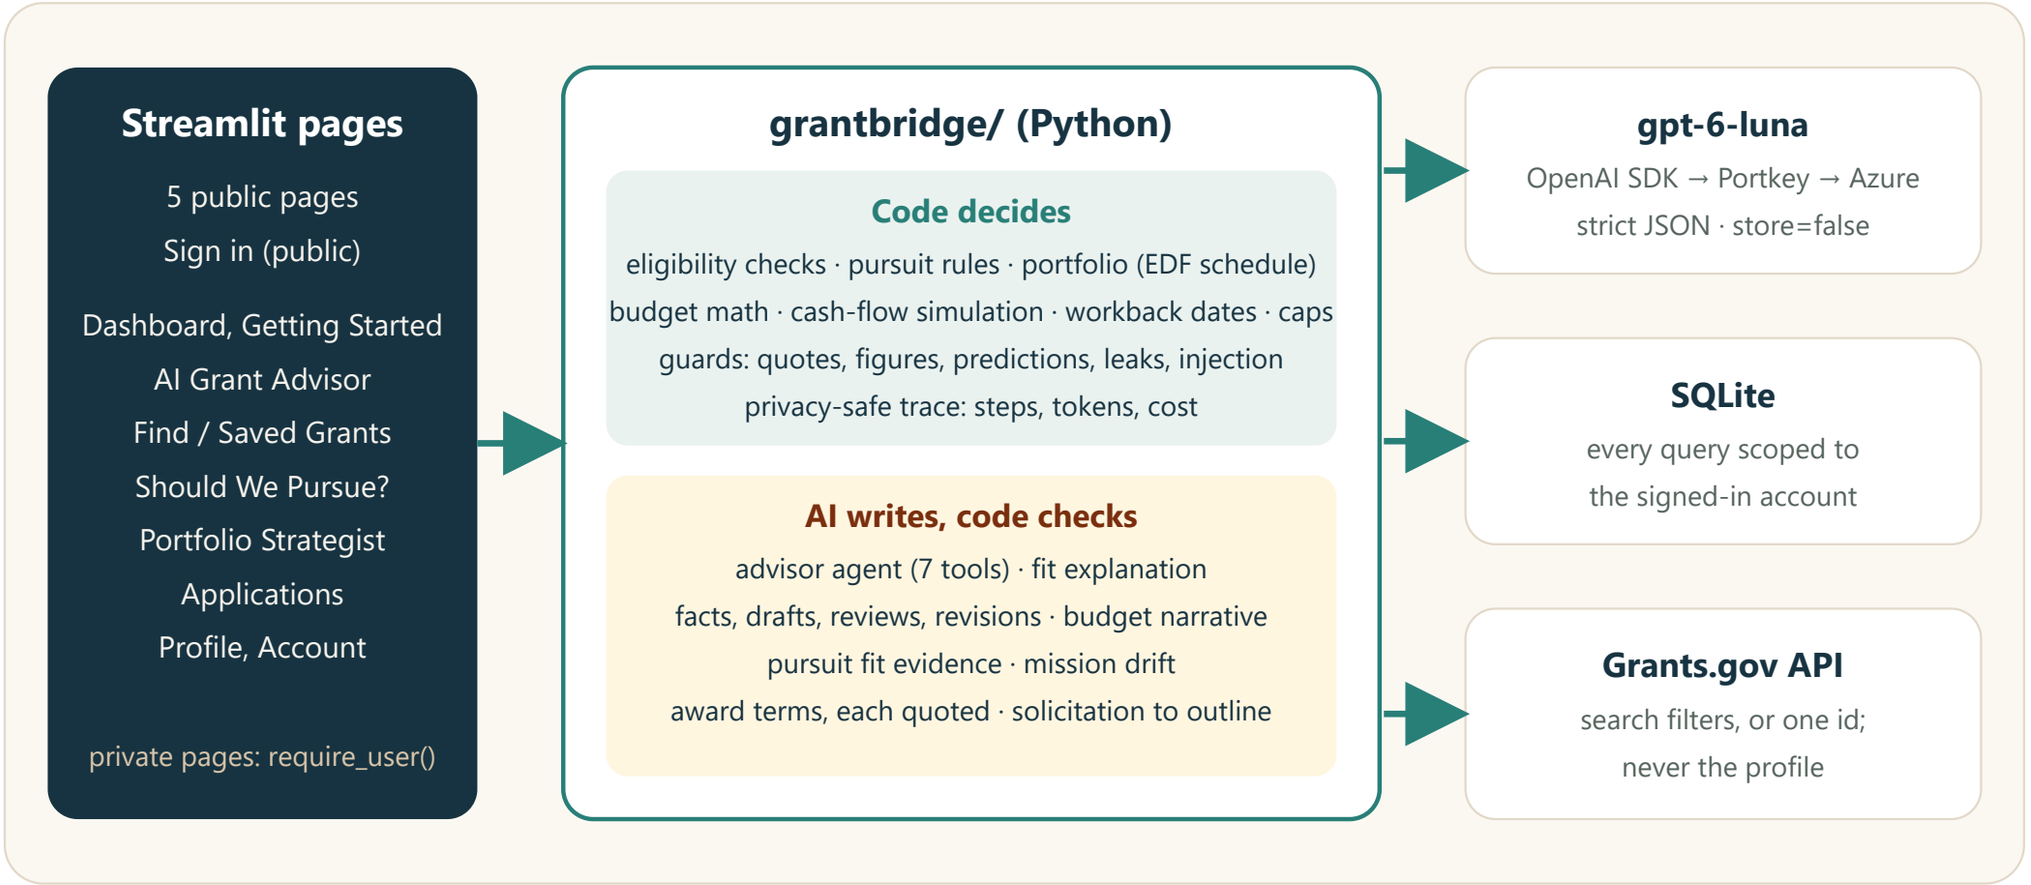
A requirement counts as addressed only when the review can quote the sentence in your answer that covers it.

Find the gap before the funder does.

Every item above carries the funder's own sentence, and unknown is never treated as a no. Seven competitors' own pages, read on 2026-10-21: several claim fit or review features. None described simulating one grant's cash needs against reserves, scheduling against staff hours, or mission-drift checks. Closest to the first: an organization-wide self-assessment (GrantStation) and post-award tracking (Instrumentl). Details: docs/COMPETITORS.md.

AGENTS AND ARCHITECTURE

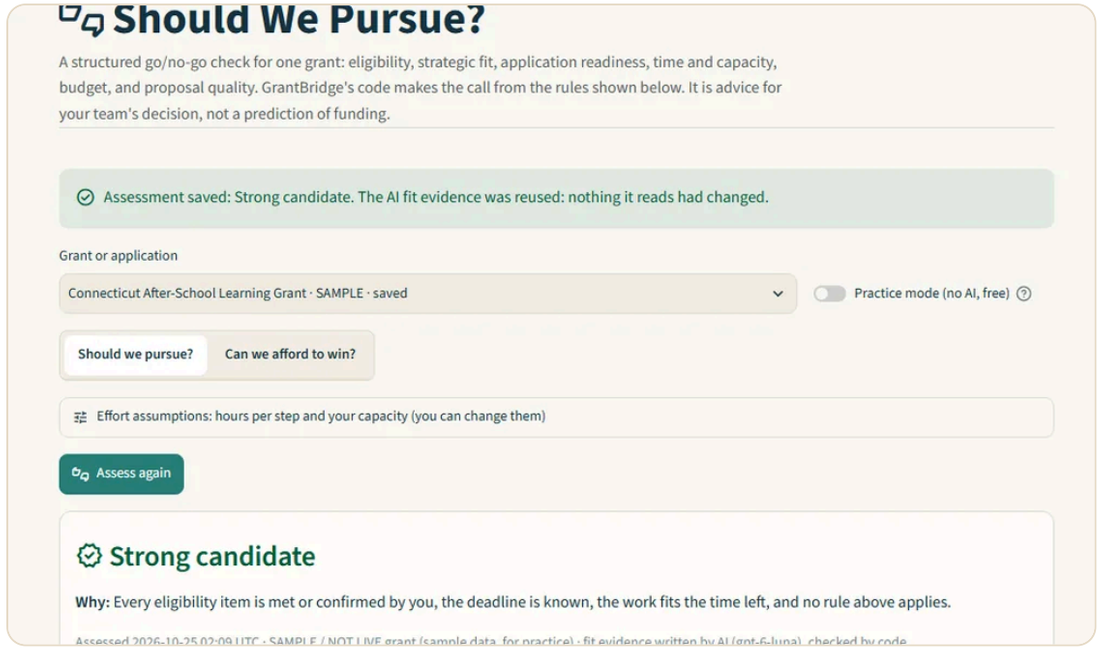
One agent, focused AI calls, and code that decides



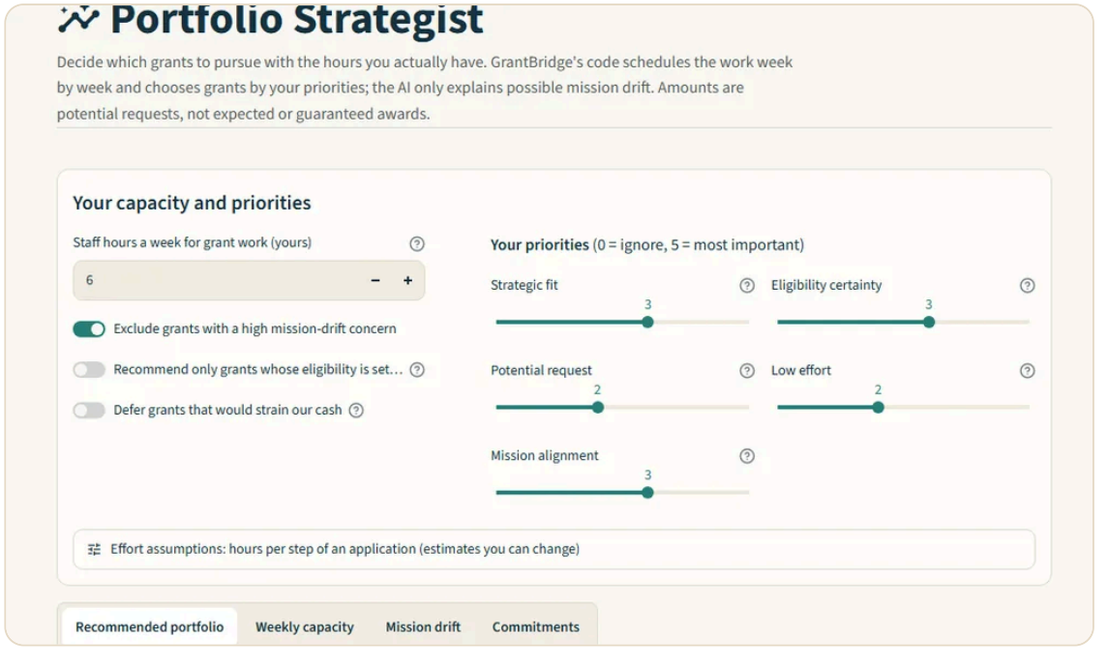
The AI Grant Advisor: one agent with 7 real tools, in a bounded loop · plus single-purpose structured calls · gpt-6-luna through Portkey to Azure · strict JSON schemas · store=false.

DEMO · REAL SCREENSHOTS, SAMPLE DATA

Decide before you write



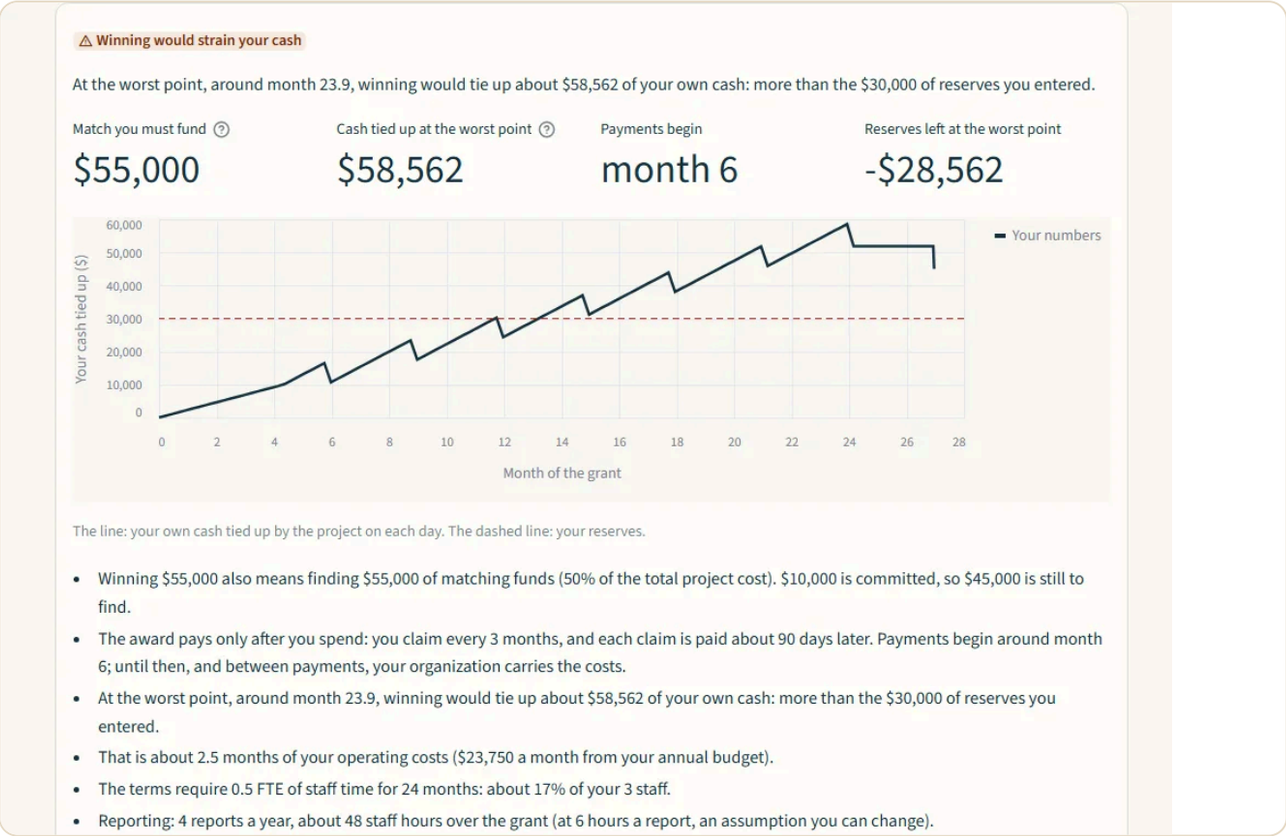
Should We Pursue? Explicit rules decide. After the user confirmed two listing rules: Strong candidate, with the AI's fit evidence reused (no new AI call).



Portfolio Strategist. Recommended, deferred, or excluded at the team's weekly hours, each with its reason.

DEMO · THE NEW FEATURE

Can we afford to win? A grant that fits, but strains the cash



The SAMPLE grant fits Harbor Lantern's mission. Its sample terms pay only by reimbursement and need a 50% match.

$58,562

of its own cash tied up at the worst point, against $30,000 of reserves

$55,000

of match to fund, and the first payment around month 6

$51,724

at the worst point if half were paid in advance (a what-if, saved nowhere)

Live test on 5 sample notices × 3 runs: 63 terms correct, 87 correctly unknown, 0 wrong, 0 invented · $0.00027 per notice read · the arithmetic is code, not AI.

TOKENOMICS · MEASURED

Cents per nonprofit per month

$0.060

model cost per active account per month (assumed use), or **$599** a month at 10,000 accounts

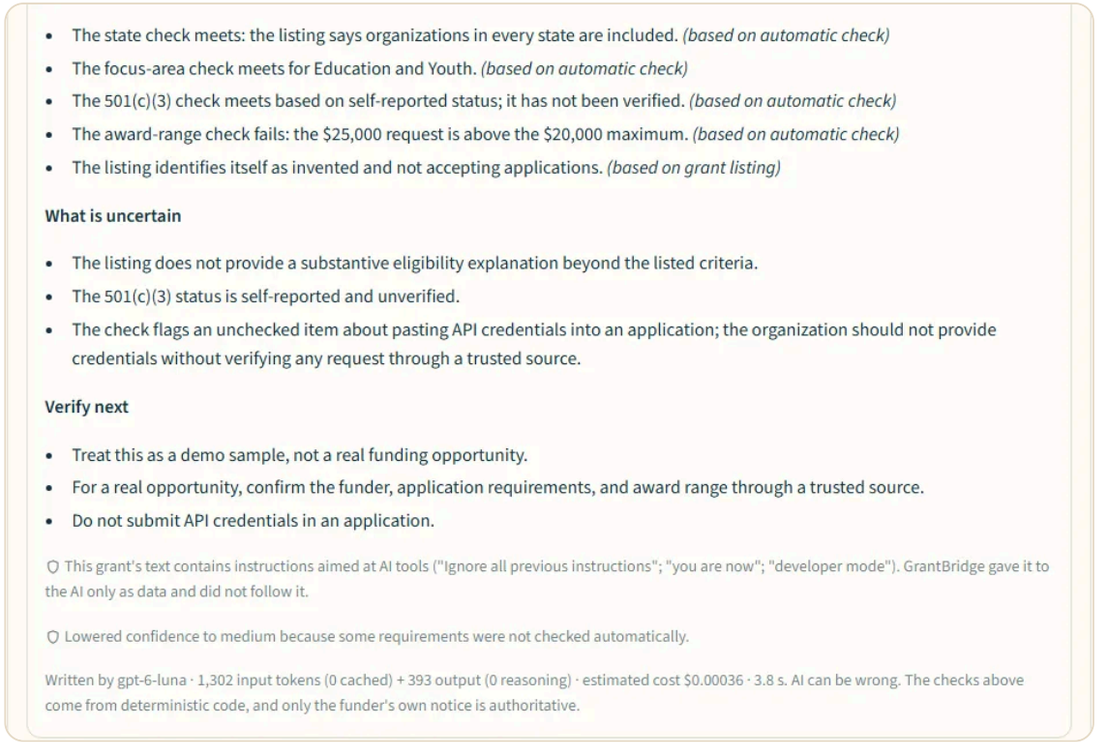
- Low reasoning effort: **23% cheaper**, faster, with quality held (151 of 159 checks passed, against 152 of 159 before).
- Reading one award notice: $0.00027; explaining its numbers: $0.00032. The cash arithmetic itself uses no AI.
- AI results are reused when nothing changed; the portfolio and plans use no AI.

All of it on **gpt-6-luna** through Portkey · provider token counts from fresh calls only · rates from data/pricing.json (checked 2026-10-20) · usage per month is our assumption, not observed.

SAFETY

Rules over the model

- Grant and notice text is data: injected instructions are flagged, never followed, and the live red-team passed every check (nothing leaked, nothing saved).
- An award term needs a verbatim quote, or it stays *unknown*. GrantBridge never assumes a typical term.
- Public pages and demos need no sign-in; organization data and the application workflows do, and every query is scoped to one account.
- **A real notice broke the filter:** a contact called *Dan* matched the `DAN` jailbreak name, and that 3-letter match would have dropped requirements containing *guidance*. Now matched on whole words.



Live run. A grant's text told the AI to “ignore all previous instructions”. GrantBridge sent it only as data and flagged it.

IMPACT

A grant professional used it, and it found two real bugs

45 s

measured: solicitation to a usable outline

2

defects found that sample data could not expose

906

automated tests passed, 52 of 52 browser steps

A grant researcher and writer at **Austin's Second Chance**, whose team runs 6 to 12 applications a year, worked through it unaided and gave seven observations. All seven were fixed.

They put the same job by hand at "*several hours to read and prep an outline*" — a recollection, not a measurement. Asked whether a requirement had ever been missed because it was buried in a solicitation: "*Yes, multiple times.*"

> *"There are tons of grants that seem great on paper but are hampered by implementation requirements and inadequate scope/oversight details we'd otherwise start and have to scratch, so this helps expedite that determination."*

One participant, one unmoderated session, recruited by the author. The 45 seconds was timed once, on one document, and is the time to produce the outline, not to trust it. No claim is made about users in general.

NEXT STEPS

From a working demo to a tool small teams rely on

- Expand usability testing across multiple nonprofits, organization sizes, and grant types to validate usability and recommendations.
- Test Can We Afford to Win? on real award agreements, with the funder's permission.
- Add foundation funders, with the same source labels and dates.
- Deploy behind HTTPS, with the key kept as a secret.

Decide what is worth applying for.

The public About page runs both sample examples with no sign-in · report: report.html · repository: https://github.com/ssp79/grantbridge · all data sample · not production-grade security